KESHAV NARAM

+1 (240) 486-6826 • naramkeshav59@gmail.com • linkedin.com/in/keshav-naram-33a834285 • github.com/Naramkeshav59

Summary

ML engineer who builds ML systems end to end, from LLM agents and speech and vision models to the cloud and data infrastructure they run on. Measures every system against ground truth and states the baseline.

Skills

ML / AI	PyTorch, Hugging Face Transformers, PEFT/QLoRA, LangChain/LangGraph agents, RAG (ChromaDB, BM25, cross-encoder rerankers), Whisper / Wav2Vec 2.0, OpenCV, scikit-learn
MLOps / Cloud	AWS (Lambda, API Gateway, SQS, ECS Fargate, S3, DynamoDB, CloudFormation), Docker, Kubernetes, Airflow, MLflow, Jenkins, GitHub Actions
Languages / Tools	Python, SQL, Bash, Git, Linux, Streamlit, Flask, pytest

Projects

ShuttleCast: AI Tactical Analyst for Pro Badminton — *LangGraph, PyTorch, AWS* 2026

- Trained stroke-detection models (frozen DINOv2 + temporal CNN) on ShuttleSet (KDD '23) so unlabelled match video becomes analyzable: hit-detection F1 0.87 on raw broadcast of a held-out match, vs 0.37 for an audio baseline.
- Built a LangGraph ReAct agent that decides per rally whether to analyze and writes the coaching note; evals caught it blaming rally winners, and a prompt fix raised correct attribution from 62% to 94% (A/B, decisions held fixed).
- Deployed on AWS (API Gateway, Lambda, SQS, ECS Fargate, S3, DynamoDB) with video upload and live results; QLoRA-tuned a 2B vision-language model from 0% to 100% output-format compliance on a $0.20 GPU run.

FHIRFlow: Making Hospital Lab Data Queryable — *FHIR, DuckDB, LLMs* 2026

- Built a FHIR R4 to DuckDB pipeline over MIMIC-IV and Synthea with data-quality rules as SQL, which showed 0% of MIMIC's 107K lab results were queryable by LOINC: 498 hospital-local codes no standard query could find.
- Built a retrieval-grounded LLM layer mapping local lab codes to LOINC (picks only retrieved codes, else abstains): 65% exact match vs 17.5% for a rule baseline on a 40-code dev set, correctly coding 87% of those lab rows.

FinSight: Cited Q&A over SEC 10-K Filings — *RAG, ChromaDB, Ollama/Groq* 2026

- Built a RAG system over SEC EDGAR 10-K filings (AAPL, JPM, JNJ, FY2021–23) with hybrid dense + BM25 retrieval, optional cross-encoder reranking, and answers with numbered citations.
- Built a benchmarking harness comparing 9 chunking × retrieval strategies on answer similarity, context relevance and latency, with a Streamlit dashboard and CI (ruff) on every push.

Experience

Software Intern – AI/ML — *WyzMindz Solutions Pvt. Ltd. (Remote)* Jul 2022 – Dec 2022

- Implemented real-time voice activity detection to segment live audio streams for speech recognition.
- Fine-tuned Wav2Vec 2.0 models on noisy call-center audio and integrated Kaldi/Fairseq-based recognizers.
- Built a Flask real-time API combining VAD, speech-to-text and live transcription in a user-facing interface.

Education

University of Maryland, College Park May 2026
M.S. Applied Machine Learning GPA 3.73 / 4.00

Indian Institute of Technology, Dharwad 2020 – 2024
B.Tech. Electrical Engineering GPA 8.08 / 10.00

Certifications & Leadership

- DeepLearning.AI: Neural Networks and Deep Learning; Natural Language Processing Specialization.
- Academic mentor, IIT Dharwad (Aug 2022 – Aug 2023): guided 20+ undergraduates on course and career planning.